

pre-final

May 4, 2024

```
[2]: # Import libraries
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers
from keras.applications import MobileNetV2
from keras.preprocessing.image import ImageDataGenerator
import numpy as np
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import cv2
import pandas as pd

import cv2
import numpy as np
import pandas as pd
import seaborn as sns
sns.set_style('darkgrid')
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from sklearn.metrics import confusion_matrix, classification_report

import tensorflow as tf
from tensorflow import keras
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam, Adamax
from tensorflow.keras.metrics import categorical_crossentropy
from tensorflow.keras.preprocessing.image import ImageDataGenerator
from tensorflow.keras.layers import Conv2D, MaxPooling2D, Flatten, Dense,
↳Activation, Dropout, BatchNormalization
from tensorflow.keras import regularizers

# Ignore Warnings
import warnings
warnings.filterwarnings("ignore")

import os
import time
```

```
import shutil
import pathlib
import itertools
```

```
/opt/conda/lib/python3.10/site-packages/scipy/__init__.py:146: UserWarning: A
NumPy version >=1.16.5 and <1.23.0 is required for this version of SciPy
(detected version 1.23.5
  warnings.warn(f"A NumPy version >={np_minversion} and <{np_maxversion}")
```

```
[3]: # Generate data paths with labels
def define_paths(data_dir):
    filepaths = []
    labels = []

    folds = os.listdir(data_dir)
    for fold in folds:
        foldpath = os.path.join(data_dir, fold)
        # check the folders from main directory. If there are another files,
        ↪ ignore them
        if pathlib.Path(foldpath).suffix != '':
            continue

        filelist = os.listdir(foldpath)
        for file in filelist:
            fpath = os.path.join(foldpath, file)

            # check if there are another folders
            if pathlib.Path(foldpath).suffix == '':
                # check unneeded masks
                if pathlib.Path(fpath).parts[-1] == 'masks' or pathlib.
                ↪ Path(fpath).parts[-1] == 'Masks' or pathlib.Path(fpath).parts[-1] == 'MASKS':
                    continue

            else:
                o_file = os.listdir(fpath)
                for f in o_file:
                    ipath = os.path.join(fpath, f)
                    filepaths.append(ipath)
                    labels.append(fold)

        else:
            filepaths.append(fpath)
            labels.append(fold)

    return filepaths, labels
```

```

# Concatenate data paths with labels into one dataframe ( to later be fitted
↳ into the model )
def define_df(files, classes):
    Fseries = pd.Series(files, name= 'filepaths')
    Lseries = pd.Series(classes, name='labels')
    return pd.concat([Fseries, Lseries], axis= 1)

# Split dataframe to train, valid, and test
def split_data(data_dir):
    # train dataframe
    files, classes = define_paths(data_dir)
    df = define_df(files, classes)
    strat = df['labels']
    train_df, dummy_df = train_test_split(df, train_size= 0.8, shuffle= True,
↳ random_state= 123, stratify= strat)

    # valid and test dataframe
    strat = dummy_df['labels']
    valid_df, test_df = train_test_split(dummy_df, train_size= 0.5, shuffle=
↳ True, random_state= 123, stratify= strat)

    return train_df, valid_df, test_df

```

```

[4]: def create_gens (train_df, valid_df, test_df, batch_size):
    '''
    This function takes train, validation, and test dataframe and fit them into
    ↳ image data generator, because model takes data from image data generator.
    Image data generator converts images into tensors. '''

    # define model parameters
    img_size = (224, 224)
    channels = 3 # either BGR or Grayscale
    color = 'rgb'
    img_shape = (img_size[0], img_size[1], channels)

    # Recommended : use custom function for test data batch size, else we can
    ↳ use normal batch size.
    ts_length = len(test_df)
    test_batch_size = max(sorted([ts_length // n for n in range(1, ts_length +
↳ 1) if ts_length%n == 0 and ts_length/n <= 80]))
    test_steps = ts_length // test_batch_size

    # This function which will be used in image data generator for data
    ↳ augmentation, it just take the image and return it again.
    def scalar(img):

```

```

        return img

    tr_gen = ImageDataGenerator(preprocessing_function= scalar,
    ↪horizontal_flip= True)
    ts_gen = ImageDataGenerator(preprocessing_function= scalar)

    train_gen = tr_gen.flow_from_dataframe( train_df, x_col= 'filepaths',
    ↪y_col= 'labels', target_size= img_size, class_mode= 'categorical',
                                         color_mode= color, shuffle= True,
    ↪batch_size= batch_size)

    valid_gen = ts_gen.flow_from_dataframe( valid_df, x_col= 'filepaths',
    ↪y_col= 'labels', target_size= img_size, class_mode= 'categorical',
                                         color_mode= color, shuffle= True,
    ↪batch_size= batch_size)

    # Note: we will use custom test_batch_size, and make shuffle= false
    test_gen = ts_gen.flow_from_dataframe( test_df, x_col= 'filepaths', y_col=
    ↪'labels', target_size= img_size, class_mode= 'categorical',
                                         color_mode= color, shuffle= False,
    ↪batch_size= test_batch_size)

    return train_gen, valid_gen, test_gen

```

```

[5]: def plot_training(hist):
    """
    This function take training model and plot history of accuracy and losses,
    ↪with the best epoch in both of them.
    """

    # Define needed variables
    tr_acc = hist.history['accuracy']
    tr_loss = hist.history['loss']
    val_acc = hist.history['val_accuracy']
    val_loss = hist.history['val_loss']
    index_loss = np.argmin(val_loss)
    val_lowest = val_loss[index_loss]
    index_acc = np.argmax(val_acc)
    acc_highest = val_acc[index_acc]
    Epochs = [i+1 for i in range(len(tr_acc))]
    loss_label = f'best epoch= {str(index_loss + 1)}'
    acc_label = f'best epoch= {str(index_acc + 1)}'

    # Plot training history
    plt.figure(figsize= (20, 8))
    plt.style.use('fivethirtyeight')

```

```

plt.subplot(1, 2, 1)
plt.plot(Epochs, tr_loss, 'r', label= 'Training loss')
plt.plot(Epochs, val_loss, 'g', label= 'Validation loss')
plt.scatter(index_loss + 1, val_lowest, s= 150, c= 'blue', label=
↳loss_label)
plt.title('Training and Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()

plt.subplot(1, 2, 2)
plt.plot(Epochs, tr_acc, 'r', label= 'Training Accuracy')
plt.plot(Epochs, val_acc, 'g', label= 'Validation Accuracy')
plt.scatter(index_acc + 1, acc_highest, s= 150, c= 'blue', label=
↳acc_label)
plt.title('Training and Validation Accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()

plt.tight_layout
plt.show()

```

```

[35]: def plot_confusion_matrix(cm, classes, normalize= False, title= 'Confusion
↳Matrix', cmap= plt.cm.Greens):
    '''
    This function plot confusion matrix method from sklearn package.
    '''

    plt.figure(figsize= (10, 10))
    plt.imshow(cm, interpolation= 'nearest', cmap= cmap)
    plt.title(title)
    plt.colorbar()

    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes, rotation= 45)
    plt.yticks(tick_marks, classes)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis= 1)[:, np.newaxis]
        print('Normalized Confusion Matrix')

    else:
        print('Confusion Matrix, Without Normalization')

    print(cm)

```

```

        thresh = cm.max() / 2.
        for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
            plt.text(j, i, cm[i, j], horizontalalignment= 'center', color=
↪'white' if cm[i, j] > thresh else 'black')

plt.tight_layout()
plt.ylabel('True Label')
plt.xlabel('Predicted Label')

```

```

[7]: def show_images(gen):
    '''
    This function take the data generator and show sample of the images
    '''

    # return classes , images to be displayed
    g_dict = gen.class_indices          # defines dictionary {'class': index}
    classes = list(g_dict.keys())       # defines list of dictionary's keys
    ↪(classes), classes names : string
    images, labels = next(gen)         # get a batch size samples from the
    ↪generator

    # calculate number of displayed samples
    length = len(labels)                # length of batch size
    sample = min(length, 25)           # check if sample less than 25 images

    plt.figure(figsize= (20, 20))

    for i in range(sample):
        plt.subplot(5, 5, i + 1)
        image = images[i] / 255        # scales data to range (0 - 255)
        plt.imshow(image)
        index = np.argmax(labels[i])   # get image index
        class_name = classes[index]    # get class of image
        plt.title(class_name, color= 'blue', fontsize= 12)
        plt.axis('off')
    plt.show()

```

```

[8]: class MyCallback(keras.callbacks.Callback):
    def __init__(self, model, patience, stop_patience, threshold, factor,
    ↪batches, epochs, ask_epoch):
        super(MyCallback, self).__init__()
        self.model = model

        self.patience = patience # specifies how many epochs without
    ↪improvement before learning rate is adjusted

```

```

        self.stop_patience = stop_patience # specifies how many times to adjust
↳lr without improvement to stop training
        self.threshold = threshold # specifies training accuracy threshold when
↳lr will be adjusted based on validation loss
        self.factor = factor # factor by which to reduce the learning rate
        self.batches = batches # number of training batch to run per epoch
        self.epochs = epochs
        self.ask_epoch = ask_epoch
        self.ask_epoch_initial = ask_epoch # save this value to restore if
↳restarting training

        # callback variables
        self.count = 0 # how many times lr has been reduced without improvement
        self.stop_count = 0
        self.best_epoch = 1 # epoch with the lowest loss
        self.initial_lr = float(tf.keras.backend.get_value(model.optimizer.lr))
↳# get the initial learning rate and save it
        self.highest_tracc = 0.0 # set highest training accuracy to 0 initially
        self.lowest_vloss = np.inf # set lowest validation loss to infinity
↳initially
        self.best_weights = self.model.get_weights() # set best weights to
↳model's initial weights
        self.initial_weights = self.model.get_weights() # save initial
↳weights if they have to get restored

        # Define a function that will run when train begins
        def on_train_begin(self, logs= None):
            msg = 'Do you want model asks you to halt the training [y/n] ?'
            print(msg)
            ans = "n" #input('')
            if ans in ['Y', 'y']:
                self.ask_permission = 1
            elif ans in ['N', 'n']:
                self.ask_permission = 0

            msg = '{0:~8s}{1:~10s}{2:~9s}{3:~9s}{4:~9s}{5:~9s}{6:~9s}{7:~10s}{8:~10s}{9:~8s}'.format('Epoch', 'Loss', 'Accuracy', 'V_loss', 'V_acc', 'LR',
↳'Next LR', 'Monitor', '% Improv', 'Duration')
            print(msg)
            self.start_time = time.time()

        def on_train_end(self, logs= None):
            stop_time = time.time()
            tr_duration = stop_time - self.start_time
            hours = tr_duration // 3600

```

```

minutes = (tr_duration - (hours * 3600)) // 60
seconds = tr_duration - ((hours * 3600) + (minutes * 60))

msg = f'training elapsed time was {str(hours)} hours, {minutes:4.1f}
↳minutes, {seconds:4.2f} seconds)'
print(msg)

# set the weights of the model to the best weights
self.model.set_weights(self.best_weights)

def on_train_batch_end(self, batch, logs= None):
    # get batch accuracy and loss
    acc = logs.get('accuracy') * 100
    loss = logs.get('loss')

    # prints over on the same line to show running batch count
    msg = '{0:20s}processing batch {1:} of {2:5s}- accuracy= {3:5.3f}
↳loss: {4:8.5f}'.format(' ', str(batch), str(self.batches), acc, loss)
    print(msg, '\r', end= '')

def on_epoch_begin(self, epoch, logs= None):
    self.ep_start = time.time()

# Define method runs on the end of each epoch
def on_epoch_end(self, epoch, logs= None):
    ep_end = time.time()
    duration = ep_end - self.ep_start

    lr = float(tf.keras.backend.get_value(self.model.optimizer.lr)) # get
↳the current learning rate
    current_lr = lr
    acc = logs.get('accuracy') # get training accuracy
    v_acc = logs.get('val_accuracy') # get validation accuracy
    loss = logs.get('loss') # get training loss for this epoch
    v_loss = logs.get('val_loss') # get the validation loss for this epoch

    if acc < self.threshold: # if training accuracy is below threshold
↳adjust lr based on training accuracy
        monitor = 'accuracy'
        if epoch == 0:
            pimprov = 0.0
        else:
            pimprov = (acc - self.highest_tracc ) * 100 / self.
↳highest_tracc # define improvement of model progres

```



```

        if acc > self.highest_tracc: # training accuracy improved in the
↪epoch
            self.highest_tracc = acc # set new highest training accuracy
            self.best_weights = self.model.get_weights() # training
↪accuracy improved so save the weights
            self.count = 0 # set count to 0 since training accuracy improved
            self.stop_count = 0 # set stop counter to 0
            if v_loss < self.lowest_vloss:
                self.lowest_vloss = v_loss
            self.best_epoch = epoch + 1 # set the value of best epoch for
↪this epoch

        else:
            # training accuracy did not improve check if this has happened
↪for patience number of epochs
            # if so adjust learning rate
            if self.count >= self.patience - 1: # lr should be adjusted
                lr = lr * self.factor # adjust the learning by factor
                tf.keras.backend.set_value(self.model.optimizer.lr, lr) #
↪set the learning rate in the optimizer
                self.count = 0 # reset the count to 0
                self.stop_count = self.stop_count + 1 # count the number of
↪consecutive lr adjustments
                self.count = 0 # reset counter
                if v_loss < self.lowest_vloss:
                    self.lowest_vloss = v_loss
                else:
                    self.count = self.count + 1 # increment patience counter

            else: # training accuracy is above threshold so adjust learning rate
↪based on validation loss
                monitor = 'val_loss'
                if epoch == 0:
                    pimprov = 0.0

                else:
                    pimprov = (self.lowest_vloss - v_loss ) * 100 / self.
↪lowest_vloss

                if v_loss < self.lowest_vloss: # check if the validation loss
↪improved
                    self.lowest_vloss = v_loss # replace lowest validation loss
↪with new validation loss
                    self.best_weights = self.model.get_weights() # validation loss
↪improved so save the weights

```

```

        self.count = 0 # reset count since validation loss improved
        self.stop_count = 0
        self.best_epoch = epoch + 1 # set the value of the best epoch
↳to this epoch

    else: # validation loss did not improve
        if self.count >= self.patience - 1: # need to adjust lr
            lr = lr * self.factor # adjust the learning rate
            self.stop_count = self.stop_count + 1 # increment stop
↳counter because lr was adjusted
            self.count = 0 # reset counter
            tf.keras.backend.set_value(self.model.optimizer.lr, lr) #
↳set the learning rate in the optimizer

        else:
            self.count = self.count + 1 # increment the patience counter

        if acc > self.highest_tracc:
            self.highest_tracc = acc

        msg = f'{str(epoch + 1):^3s}/{str(self.epochs):4s} {loss:^9.3f}{acc *
↳100:^9.3f}{v_loss:^9.5f}{v_acc * 100:^9.3f}{current_lr:^9.5f}{lr:^9.
↳5f}{monitor:^11s}{pimprov:^10.2f}{duration:^8.2f}'
        print(msg)

        if self.stop_count > self.stop_patience - 1: # check if learning rate
↳has been adjusted stop_count times with no improvement
            msg = f' training has been halted at epoch {epoch + 1} after {self.
↳stop_patience} adjustments of learning rate with no improvement'
            print(msg)
            self.model.stop_training = True # stop training

    else:
        if self.ask_epoch != None and self.ask_permission != 0:
            if epoch + 1 >= self.ask_epoch:
                msg = 'enter H to halt training or an integer for number of
↳epochs to run then ask again'
                print(msg)

                ans = input('')
                if ans == 'H' or ans == 'h':
                    msg = f'training has been halted at epoch {epoch + 1}
↳due to user input'
                    print(msg)
                    self.model.stop_training = True # stop training

```

```

        else:
            try:
                ans = int(ans)
                self.ask_epoch += ans
                msg = f' training will continue until epoch_
↳{str(self.ask_epoch)}'

                print(msg)
                msg = '{0:^8s}{1:^10s}{2:^9s}{3:^9s}{4:^9s}{5:
↳^9s}{6:^9s}{7:^10s}{8:10s}{9:^8s}'.format('Epoch', 'Loss', 'Accuracy',
↳'V_loss', 'V_acc', 'LR', 'Next LR', 'Monitor', '% Improv', 'Duration')
                print(msg)

            except Exception:
                print('Invalid')

```

```

[9]: data_dir = '/kaggle/input/covid19-radiography-database/
↳COVID-19_Radiography_Dataset'

try:
    # Get splitted data
    train_df, valid_df, test_df = split_data(data_dir)

    # Get Generators
    batch_size = 16
    train_gen, valid_gen, test_gen = create_gens(train_df, valid_df, test_df,
↳batch_size)

except:
    print('Invalid Input')

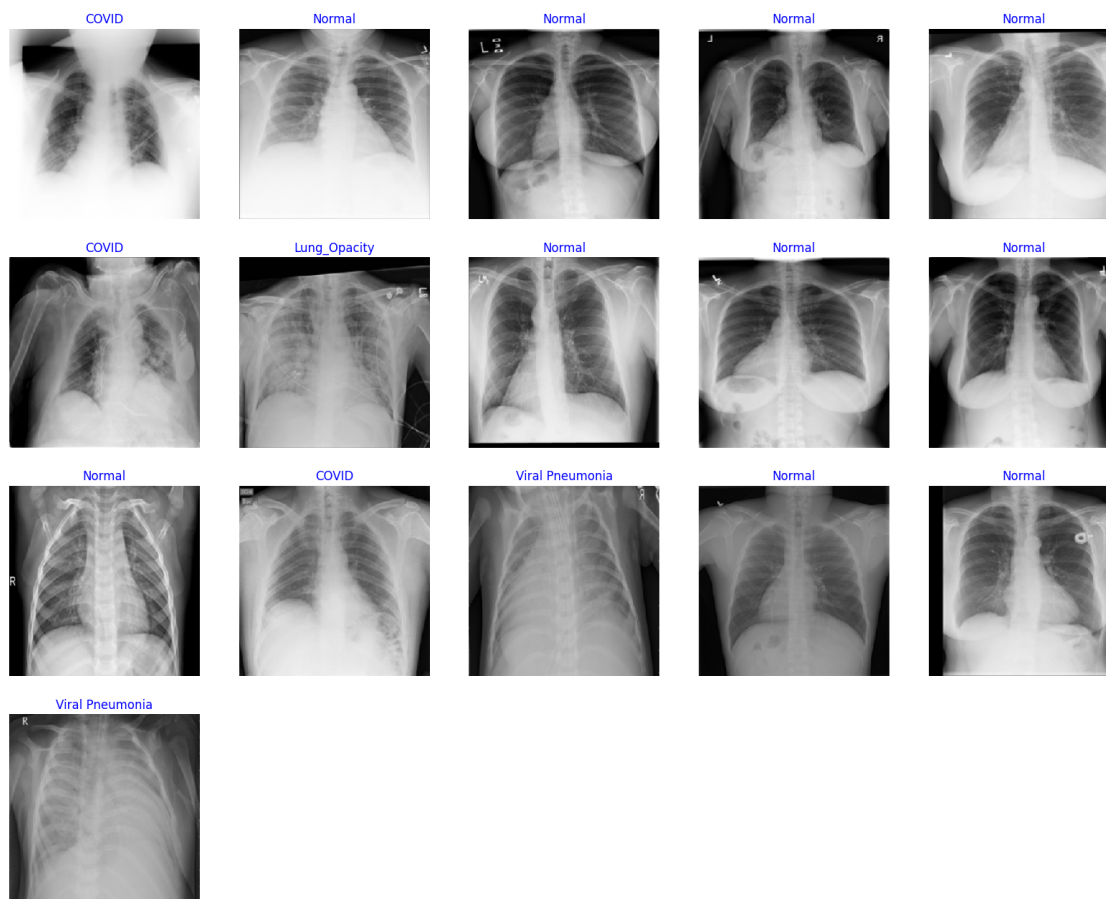
```

Found 16932 validated image filenames belonging to 4 classes.
Found 2116 validated image filenames belonging to 4 classes.
Found 2117 validated image filenames belonging to 4 classes.

```

[10]: show_images(train_gen)

```



[11]: #3 & 4. Load Pre-trained MobileNetV2

```
img_size = (224, 224)
channels = 3
img_shape = (img_size[0], img_size[1], channels)

def getBaseModel():
    base_model = MobileNetV2(
        weights="imagenet", include_top=False, input_shape= img_shape
    )
    return base_model
base_model = getBaseModel()
base_model.summary()
```

Downloading data from https://storage.googleapis.com/tensorflow/keras-applications/mobilenet_v2/mobilenet_v2_weights_tf_dim_ordering_tf_kernels_1.0_224_no_top.h5

9406464/9406464 [=====] - 0s 0us/step

Model: "mobilenetv2_1.00_224"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_1 (InputLayer)	[(None, 224, 224, 3 0)]		[]
Conv1 (Conv2D) ['input_1[0][0]']	(None, 112, 112, 32 864)		
bn_Conv1 (BatchNormalization)	(None, 112, 112, 32 128)		['Conv1[0][0]']
Conv1_relu (ReLU) ['bn_Conv1[0][0]']	(None, 112, 112, 32 0)		
expanded_conv_depthwise (Depth ['Conv1_relu[0][0]'] wiseConv2D)	(None, 112, 112, 32 288)		
expanded_conv_depthwise_BN (Ba ['expanded_conv_depthwise[0][0]'] tchNormalization)	(None, 112, 112, 32 128)		
expanded_conv_depthwise_relu (['expanded_conv_depthwise_BN[0][0] ReLU)	(None, 112, 112, 32 0)		['']
expanded_conv_project (Conv2D) ['expanded_conv_depthwise_relu[0])	(None, 112, 112, 16 512)		[0]']
expanded_conv_project_BN (Batc ['expanded_conv_project[0][0]'] hNormalization)	(None, 112, 112, 16 64)		
block_1_expand (Conv2D) ['expanded_conv_project_BN[0][0]'])	(None, 112, 112, 96 1536)]
block_1_expand_BN (BatchNormal ['block_1_expand[0][0]'] ization)	(None, 112, 112, 96 384)		
block_1_expand_relu (ReLU)	(None, 112, 112, 96 0		

```

['block_1_expand_BN[0][0]']
)

block_1_pad (ZeroPadding2D) (None, 113, 113, 96) 0
['block_1_expand_relu[0][0]']
)

block_1_depthwise (DepthwiseCo (None, 56, 56, 96) 864
['block_1_pad[0][0]']
nv2D)

block_1_depthwise_BN (BatchNor (None, 56, 56, 96) 384
['block_1_depthwise[0][0]']
malization)

block_1_depthwise_relu (ReLU) (None, 56, 56, 96) 0
['block_1_depthwise_BN[0][0]']

block_1_project (Conv2D) (None, 56, 56, 24) 2304
['block_1_depthwise_relu[0][0]']

block_1_project_BN (BatchNorma (None, 56, 56, 24) 96
['block_1_project[0][0]']
lization)

block_2_expand (Conv2D) (None, 56, 56, 144) 3456
['block_1_project_BN[0][0]']

block_2_expand_BN (BatchNormal (None, 56, 56, 144) 576
['block_2_expand[0][0]']
ization)

block_2_expand_relu (ReLU) (None, 56, 56, 144) 0
['block_2_expand_BN[0][0]']

block_2_depthwise (DepthwiseCo (None, 56, 56, 144) 1296
['block_2_expand_relu[0][0]']
nv2D)

block_2_depthwise_BN (BatchNor (None, 56, 56, 144) 576
['block_2_depthwise[0][0]']
malization)

block_2_depthwise_relu (ReLU) (None, 56, 56, 144) 0
['block_2_depthwise_BN[0][0]']

block_2_project (Conv2D) (None, 56, 56, 24) 3456
['block_2_depthwise_relu[0][0]']

```

```

block_2_project_BN (BatchNorma (None, 56, 56, 24) 96
['block_2_project[0][0]']
lization)

block_2_add (Add) (None, 56, 56, 24) 0
['block_1_project_BN[0][0]',
'block_2_project_BN[0][0]']

block_3_expand (Conv2D) (None, 56, 56, 144) 3456
['block_2_add[0][0]']

block_3_expand_BN (BatchNormal (None, 56, 56, 144) 576
['block_3_expand[0][0]']
ization)

block_3_expand_relu (ReLU) (None, 56, 56, 144) 0
['block_3_expand_BN[0][0]']

block_3_pad (ZeroPadding2D) (None, 57, 57, 144) 0
['block_3_expand_relu[0][0]']

block_3_depthwise (DepthwiseCo (None, 28, 28, 144) 1296
['block_3_pad[0][0]']
nv2D)

block_3_depthwise_BN (BatchNor (None, 28, 28, 144) 576
['block_3_depthwise[0][0]']
malization)

block_3_depthwise_relu (ReLU) (None, 28, 28, 144) 0
['block_3_depthwise_BN[0][0]']

block_3_project (Conv2D) (None, 28, 28, 32) 4608
['block_3_depthwise_relu[0][0]']

block_3_project_BN (BatchNorma (None, 28, 28, 32) 128
['block_3_project[0][0]']
lization)

block_4_expand (Conv2D) (None, 28, 28, 192) 6144
['block_3_project_BN[0][0]']

block_4_expand_BN (BatchNormal (None, 28, 28, 192) 768
['block_4_expand[0][0]']
ization)

block_4_expand_relu (ReLU) (None, 28, 28, 192) 0

```

```

['block_4_expand_BN[0][0]']

block_4_depthwise (DepthwiseCo (None, 28, 28, 192) 1728
['block_4_expand_relu[0][0]']
nv2D)

block_4_depthwise_BN (BatchNor (None, 28, 28, 192) 768
['block_4_depthwise[0][0]']
malization)

block_4_depthwise_relu (ReLU) (None, 28, 28, 192) 0
['block_4_depthwise_BN[0][0]']

block_4_project (Conv2D) (None, 28, 28, 32) 6144
['block_4_depthwise_relu[0][0]']

block_4_project_BN (BatchNorma (None, 28, 28, 32) 128
['block_4_project[0][0]']
lization)

block_4_add (Add) (None, 28, 28, 32) 0
['block_3_project_BN[0][0]',
'block_4_project_BN[0][0]']

block_5_expand (Conv2D) (None, 28, 28, 192) 6144
['block_4_add[0][0]']

block_5_expand_BN (BatchNormal (None, 28, 28, 192) 768
['block_5_expand[0][0]']
ization)

block_5_expand_relu (ReLU) (None, 28, 28, 192) 0
['block_5_expand_BN[0][0]']

block_5_depthwise (DepthwiseCo (None, 28, 28, 192) 1728
['block_5_expand_relu[0][0]']
nv2D)

block_5_depthwise_BN (BatchNor (None, 28, 28, 192) 768
['block_5_depthwise[0][0]']
malization)

block_5_depthwise_relu (ReLU) (None, 28, 28, 192) 0
['block_5_depthwise_BN[0][0]']

block_5_project (Conv2D) (None, 28, 28, 32) 6144
['block_5_depthwise_relu[0][0]']

```



```

block_5_project_BN (BatchNorma (None, 28, 28, 32) 128
['block_5_project[0][0]']
lization)

block_5_add (Add) (None, 28, 28, 32) 0
['block_4_add[0][0]',
'block_5_project_BN[0][0]']

block_6_expand (Conv2D) (None, 28, 28, 192) 6144
['block_5_add[0][0]']

block_6_expand_BN (BatchNormal (None, 28, 28, 192) 768
['block_6_expand[0][0]']
ization)

block_6_expand_relu (ReLU) (None, 28, 28, 192) 0
['block_6_expand_BN[0][0]']

block_6_pad (ZeroPadding2D) (None, 29, 29, 192) 0
['block_6_expand_relu[0][0]']

block_6_depthwise (DepthwiseCo (None, 14, 14, 192) 1728
['block_6_pad[0][0]']
nv2D)

block_6_depthwise_BN (BatchNor (None, 14, 14, 192) 768
['block_6_depthwise[0][0]']
malization)

block_6_depthwise_relu (ReLU) (None, 14, 14, 192) 0
['block_6_depthwise_BN[0][0]']

block_6_project (Conv2D) (None, 14, 14, 64) 12288
['block_6_depthwise_relu[0][0]']

block_6_project_BN (BatchNorma (None, 14, 14, 64) 256
['block_6_project[0][0]']
lization)

block_7_expand (Conv2D) (None, 14, 14, 384) 24576
['block_6_project_BN[0][0]']

block_7_expand_BN (BatchNormal (None, 14, 14, 384) 1536
['block_7_expand[0][0]']
ization)

block_7_expand_relu (ReLU) (None, 14, 14, 384) 0
['block_7_expand_BN[0][0]']

```

```

block_7_depthwise (DepthwiseCo (None, 14, 14, 384) 3456
['block_7_expand_relu[0][0]']
nv2D)

block_7_depthwise_BN (BatchNor (None, 14, 14, 384) 1536
['block_7_depthwise[0][0]']
malization)

block_7_depthwise_relu (ReLU) (None, 14, 14, 384) 0
['block_7_depthwise_BN[0][0]']

block_7_project (Conv2D) (None, 14, 14, 64) 24576
['block_7_depthwise_relu[0][0]']

block_7_project_BN (BatchNorma (None, 14, 14, 64) 256
['block_7_project[0][0]']
lization)

block_7_add (Add) (None, 14, 14, 64) 0
['block_6_project_BN[0][0]',
'block_7_project_BN[0][0]']

block_8_expand (Conv2D) (None, 14, 14, 384) 24576
['block_7_add[0][0]']

block_8_expand_BN (BatchNormal (None, 14, 14, 384) 1536
['block_8_expand[0][0]']
ization)

block_8_expand_relu (ReLU) (None, 14, 14, 384) 0
['block_8_expand_BN[0][0]']

block_8_depthwise (DepthwiseCo (None, 14, 14, 384) 3456
['block_8_expand_relu[0][0]']
nv2D)

block_8_depthwise_BN (BatchNor (None, 14, 14, 384) 1536
['block_8_depthwise[0][0]']
malization)

block_8_depthwise_relu (ReLU) (None, 14, 14, 384) 0
['block_8_depthwise_BN[0][0]']

block_8_project (Conv2D) (None, 14, 14, 64) 24576
['block_8_depthwise_relu[0][0]']

block_8_project_BN (BatchNorma (None, 14, 14, 64) 256

```

```

['block_8_project[0][0]']
lization)

block_8_add (Add) (None, 14, 14, 64) 0
['block_7_add[0][0]',
'block_8_project_BN[0][0]']

block_9_expand (Conv2D) (None, 14, 14, 384) 24576
['block_8_add[0][0]']

block_9_expand_BN (BatchNormal (None, 14, 14, 384) 1536
['block_9_expand[0][0]']
lization)

block_9_expand_relu (ReLU) (None, 14, 14, 384) 0
['block_9_expand_BN[0][0]']

block_9_depthwise (DepthwiseCo (None, 14, 14, 384) 3456
['block_9_expand_relu[0][0]']
nv2D)

block_9_depthwise_BN (BatchNor (None, 14, 14, 384) 1536
['block_9_depthwise[0][0]']
malization)

block_9_depthwise_relu (ReLU) (None, 14, 14, 384) 0
['block_9_depthwise_BN[0][0]']

block_9_project (Conv2D) (None, 14, 14, 64) 24576
['block_9_depthwise_relu[0][0]']

block_9_project_BN (BatchNorma (None, 14, 14, 64) 256
['block_9_project[0][0]']
lization)

block_9_add (Add) (None, 14, 14, 64) 0
['block_8_add[0][0]',
'block_9_project_BN[0][0]']

block_10_expand (Conv2D) (None, 14, 14, 384) 24576
['block_9_add[0][0]']

block_10_expand_BN (BatchNorma (None, 14, 14, 384) 1536
['block_10_expand[0][0]']
lization)

block_10_expand_relu (ReLU) (None, 14, 14, 384) 0
['block_10_expand_BN[0][0]']

```

```

block_10_depthwise (DepthwiseC (None, 14, 14, 384) 3456
['block_10_expand_relu[0][0]']
onv2D)

block_10_depthwise_BN (BatchNo (None, 14, 14, 384) 1536
['block_10_depthwise[0][0]']
rmalization)

block_10_depthwise_relu (ReLU) (None, 14, 14, 384) 0
['block_10_depthwise_BN[0][0]']

block_10_project (Conv2D) (None, 14, 14, 96) 36864
['block_10_depthwise_relu[0][0]']

block_10_project_BN (BatchNorm (None, 14, 14, 96) 384
['block_10_project[0][0]']
alization)

block_11_expand (Conv2D) (None, 14, 14, 576) 55296
['block_10_project_BN[0][0]']

block_11_expand_BN (BatchNorma (None, 14, 14, 576) 2304
['block_11_expand[0][0]']
lization)

block_11_expand_relu (ReLU) (None, 14, 14, 576) 0
['block_11_expand_BN[0][0]']

block_11_depthwise (DepthwiseC (None, 14, 14, 576) 5184
['block_11_expand_relu[0][0]']
onv2D)

block_11_depthwise_BN (BatchNo (None, 14, 14, 576) 2304
['block_11_depthwise[0][0]']
rmalization)

block_11_depthwise_relu (ReLU) (None, 14, 14, 576) 0
['block_11_depthwise_BN[0][0]']

block_11_project (Conv2D) (None, 14, 14, 96) 55296
['block_11_depthwise_relu[0][0]']

block_11_project_BN (BatchNorm (None, 14, 14, 96) 384
['block_11_project[0][0]']
alization)

block_11_add (Add) (None, 14, 14, 96) 0

```

```

['block_10_project_BN[0][0]',
'block_11_project_BN[0][0]']

block_12_expand (Conv2D)          (None, 14, 14, 576) 55296
['block_11_add[0][0]']

block_12_expand_BN (BatchNorma    (None, 14, 14, 576) 2304
['block_12_expand[0][0]']
lization)

block_12_expand_relu (ReLU)       (None, 14, 14, 576) 0
['block_12_expand_BN[0][0]']

block_12_depthwise (DepthwiseC    (None, 14, 14, 576) 5184
['block_12_expand_relu[0][0]']
onv2D)

block_12_depthwise_BN (BatchNo    (None, 14, 14, 576) 2304
['block_12_depthwise[0][0]']
rmalization)

block_12_depthwise_relu (ReLU)    (None, 14, 14, 576) 0
['block_12_depthwise_BN[0][0]']

block_12_project (Conv2D)         (None, 14, 14, 96) 55296
['block_12_depthwise_relu[0][0]']

block_12_project_BN (BatchNorm    (None, 14, 14, 96) 384
['block_12_project[0][0]']
alization)

block_12_add (Add)                (None, 14, 14, 96) 0
['block_11_add[0][0]',
'block_12_project_BN[0][0]']

block_13_expand (Conv2D)          (None, 14, 14, 576) 55296
['block_12_add[0][0]']

block_13_expand_BN (BatchNorma    (None, 14, 14, 576) 2304
['block_13_expand[0][0]']
lization)

block_13_expand_relu (ReLU)       (None, 14, 14, 576) 0
['block_13_expand_BN[0][0]']

block_13_pad (ZeroPadding2D)      (None, 15, 15, 576) 0
['block_13_expand_relu[0][0]']

```

block_13_depthwise (DepthwiseC ['block_13_pad[0][0]' onv2D)	(None, 7, 7, 576)	5184
block_13_depthwise_BN (BatchNo ['block_13_depthwise[0][0]' rmalization)	(None, 7, 7, 576)	2304
block_13_depthwise_relu (ReLU) ['block_13_depthwise_BN[0][0]']	(None, 7, 7, 576)	0
block_13_project (Conv2D) ['block_13_depthwise_relu[0][0]']	(None, 7, 7, 160)	92160
block_13_project_BN (BatchNorm ['block_13_project[0][0]' alization)	(None, 7, 7, 160)	640
block_14_expand (Conv2D) ['block_13_project_BN[0][0]']	(None, 7, 7, 960)	153600
block_14_expand_BN (BatchNorma ['block_14_expand[0][0]' lization)	(None, 7, 7, 960)	3840
block_14_expand_relu (ReLU) ['block_14_expand_BN[0][0]']	(None, 7, 7, 960)	0
block_14_depthwise (DepthwiseC ['block_14_expand_relu[0][0]' onv2D)	(None, 7, 7, 960)	8640
block_14_depthwise_BN (BatchNo ['block_14_depthwise[0][0]' rmalization)	(None, 7, 7, 960)	3840
block_14_depthwise_relu (ReLU) ['block_14_depthwise_BN[0][0]']	(None, 7, 7, 960)	0
block_14_project (Conv2D) ['block_14_depthwise_relu[0][0]']	(None, 7, 7, 160)	153600
block_14_project_BN (BatchNorm ['block_14_project[0][0]' alization)	(None, 7, 7, 160)	640
block_14_add (Add) ['block_13_project_BN[0][0]',	(None, 7, 7, 160)	0

```

'block_14_project_BN[0][0]']

block_15_expand (Conv2D)      (None, 7, 7, 960)    153600
['block_14_add[0][0]']

block_15_expand_BN (BatchNorma (None, 7, 7, 960)    3840
['block_15_expand[0][0]']
lization)

block_15_expand_relu (ReLU)    (None, 7, 7, 960)    0
['block_15_expand_BN[0][0]']

block_15_depthwise (DepthwiseC (None, 7, 7, 960)    8640
['block_15_expand_relu[0][0]']
onv2D)

block_15_depthwise_BN (BatchNo (None, 7, 7, 960)    3840
['block_15_depthwise[0][0]']
rmalization)

block_15_depthwise_relu (ReLU) (None, 7, 7, 960)    0
['block_15_depthwise_BN[0][0]']

block_15_project (Conv2D)      (None, 7, 7, 160)    153600
['block_15_depthwise_relu[0][0]']

block_15_project_BN (BatchNorm (None, 7, 7, 160)    640
['block_15_project[0][0]']
alization)

block_15_add (Add)             (None, 7, 7, 160)    0
['block_14_add[0][0]',
'block_15_project_BN[0][0]']

block_16_expand (Conv2D)      (None, 7, 7, 960)    153600
['block_15_add[0][0]']

block_16_expand_BN (BatchNorma (None, 7, 7, 960)    3840
['block_16_expand[0][0]']
lization)

block_16_expand_relu (ReLU)    (None, 7, 7, 960)    0
['block_16_expand_BN[0][0]']

block_16_depthwise (DepthwiseC (None, 7, 7, 960)    8640
['block_16_expand_relu[0][0]']
onv2D)

```

```

block_16_depthwise_BN (BatchNo (None, 7, 7, 960) 3840
['block_16_depthwise[0][0]']
rmalization)

block_16_depthwise_relu (ReLU) (None, 7, 7, 960) 0
['block_16_depthwise_BN[0][0]']

block_16_project (Conv2D) (None, 7, 7, 320) 307200
['block_16_depthwise_relu[0][0]']

block_16_project_BN (BatchNorm (None, 7, 7, 320) 1280
['block_16_project[0][0]']
alization)

Conv_1 (Conv2D) (None, 7, 7, 1280) 409600
['block_16_project_BN[0][0]']

Conv_1_bn (BatchNormalization) (None, 7, 7, 1280) 5120
['Conv_1[0][0]']

out_relu (ReLU) (None, 7, 7, 1280) 0
['Conv_1_bn[0][0]']

```

```

=====
=====

```

```

Total params: 2,257,984
Trainable params: 2,223,872
Non-trainable params: 34,112

```

```

-----
-----

```

```

[12]: # 5. Replace Fully Connected Layers

# Create Model Structure
num_classes = len(list(train_gen.class_indices.keys())) # to define number of
↳ classes in dense layer

model = keras.Sequential(
    [
        base_model,
        layers.Flatten(),
        layers.Dense(1024, activation="relu"),
        layers.Dense(512, activation="relu"),
        layers.Dense(num_classes, activation="softmax"),
    ]
)

```



```
#model.compile(Adamax(learning_rate= 0.001), loss= 'categorical_crossentropy',  
↳metrics= ['accuracy'])
```

[13]: # 6. Total Parameters

```
total_params = model.count_params()  
print(f"Total Parameters: {total_params}")
```

Total Parameters: 67011140

[14]: # 7. Print Architecture

```
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
flatten (Flatten)	(None, 62720)	0
dense (Dense)	(None, 1024)	64226304
dense_1 (Dense)	(None, 512)	524800
dense_2 (Dense)	(None, 4)	2052

=====
Total params: 67,011,140
Trainable params: 66,977,028
Non-trainable params: 34,112
=====

[15]:

```
batch_size = 64    #batch size for training  
epochs = 20    # number of all epochs in training  
patience = 3    #number of epochs to wait to adjust lr if monitored value does  
↳not improve  
stop_patience = 10    # number of epochs to wait before stopping training if  
↳monitored value does not improve  
threshold = 0.9    # if train accuracy is < threshold adjust monitor accuracy,  
↳else monitor validation loss  
factor = 0.5    # factor to reduce lr by  
ask_epoch = 50    # number of epochs to run before asking if you want to halt  
↳training
```

```

batches = int(np.ceil(len(train_gen.labels) / batch_size))    # number of
↳training batch to run per epoch
print(f'Number of batches in training: {batches*64}')
```

Number of batches in training: 16960

[16]: # 8. Retrain with Transfer Learning

```

def train_model(model):

    # Compile and train model
    model.compile(
        optimizer=keras.optimizers.Adam(learning_rate=0.001),
        loss="categorical_crossentropy",
        metrics=["accuracy"],
    )
    print(model.summary())
    callbacks = [MyCallback(model= model, patience= patience, stop_patience=
↳stop_patience, threshold= threshold,
        factor= factor, batches= batches, epochs= epochs, ask_epoch=
↳ask_epoch )]

    history = model.fit(
        x= train_gen, epochs= epochs, batch_size= batch_size, callbacks=
↳callbacks,
        validation_data= valid_gen, validation_steps= None, shuffle=
↳False, verbose=0
    )
    return history

# define a function to set number of trainable layers
def set_trainable_layers(trainable_layers):
    # Freeze base model layers
    tempBASEMODEL = getBaseModel()
    for layer in tempBASEMODEL.layers:
        layer.trainable = False

    # Unfreeze specific layers
    if trainable_layers == 4:
        for layer in tempBASEMODEL.layers[-3:]:
            layer.trainable = True
            print(f"Layer {layer.name} is trainable")
    elif trainable_layers == 5:
        for layer in tempBASEMODEL.layers[-5:]:
            layer.trainable = True
            print(f"Layer {layer.name} is trainable")
```

```

model = keras.Sequential(
[
    tempBASEMODEL,
    layers.Flatten(),
    layers.Dense(1024, activation="relu"),
    layers.Dense(512, activation="relu"),
    layers.Dense(num_classes, activation="softmax"),
]
)

return model

model1 = set_trainable_layers( trainable_layers=3) # Last 3 FC layers
model2 = set_trainable_layers( trainable_layers=4) # 1 Conv + 3 FC layers
model3 = set_trainable_layers( trainable_layers=5) # 2 Conv + 3 FC layers

# Train with different sets of trainable layers
history_3 = train_model(model1) # Last 3 FC layers
history_4 = train_model(model2) # 1 Conv + 3 FC layers
history_5 = train_model(model3) # 2 Conv + 3 FC layers

# 9. Evaluation and Comparison

# (Code for calculating accuracy, recall, precision, F1-score, sensitivity,
# and plotting convergence curves would go here)

```

```

Layer Conv_1 is trainable
Layer Conv_1_bn is trainable
Layer out_relu is trainable
Layer block_16_project is trainable
Layer block_16_project_BN is trainable
Layer Conv_1 is trainable
Layer Conv_1_bn is trainable
Layer out_relu is trainable
Model: "sequential_1"

```

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
flatten_1 (Flatten)	(None, 62720)	0
dense_3 (Dense)	(None, 1024)	64226304
dense_4 (Dense)	(None, 512)	524800

dense_5 (Dense)

(None, 4)

2052

=====

Total params: 67,011,140

Trainable params: 64,753,156

Non-trainable params: 2,257,984

None

Do you want model asks you to halt the training [y/n] ?

Epoch	Loss	Accuracy	V_loss	V_acc	LR	Next LR	Monitor	%
Improv	Duration							
1 /20	1.178	76.695	0.49289	81.994	0.00100	0.00100	accuracy	
0.00	162.75							
2 /20	0.438	83.942	0.39952	84.972	0.00100	0.00100	accuracy	
9.45	64.17							
3 /20	0.374	85.920	0.48032	83.270	0.00100	0.00100	accuracy	
2.36	64.65							
4 /20	0.326	87.863	0.39263	85.917	0.00100	0.00100	accuracy	
2.26	62.55							
5 /20	0.293	89.452	0.41083	85.350	0.00100	0.00100	accuracy	
1.81	62.71							
6 /20	0.278	89.659	0.36501	87.004	0.00100	0.00100	accuracy	
0.23	64.04							
7 /20	0.253	90.621	0.39482	86.862	0.00100	0.00100	val_loss	
-8.17	61.80							
8 /20	0.240	91.206	0.35082	86.862	0.00100	0.00100	val_loss	
3.89	64.74							
9 /20	0.213	92.251	0.41785	85.775	0.00100	0.00100	val_loss	
-19.11	62.67							
10 /20	0.210	92.245	0.47447	85.113	0.00100	0.00100	val_loss	
-35.25	62.54							
11 /20	0.196	92.623	0.36808	86.437	0.00100	0.00050	val_loss	
-4.92	62.75							
12 /20	0.143	94.673	0.38946	87.760	0.00050	0.00050	val_loss	
-11.02	62.33							
13 /20	0.131	95.263	0.38550	87.807	0.00050	0.00050	val_loss	
-9.89	63.42							
14 /20	0.122	95.659	0.41437	87.098	0.00050	0.00025	val_loss	
-18.11	61.71							
15 /20	0.095	96.781	0.42801	88.138	0.00025	0.00025	val_loss	
-22.00	61.77							
16 /20	0.093	96.687	0.47304	87.335	0.00025	0.00025	val_loss	
-34.84	62.28							
17 /20	0.084	97.000	0.46536	88.138	0.00025	0.00013	val_loss	
-32.65	61.32							
18 /20	0.069	97.632	0.51962	86.815	0.00013	0.00013	val_loss	
-48.12	66.69							
19 /20	0.062	97.868	0.51418	87.713	0.00013	0.00013	val_loss	

```

-46.57    62.70
20 /20      0.065    97.903    0.52973    87.949    0.00013    0.00006    val_loss
-51.00    62.75
training elapsed time was 0.0 hours, 22.0 minutes, 43.27 seconds)
Model: "sequential_2"

```

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
flatten_2 (Flatten)	(None, 62720)	0
dense_6 (Dense)	(None, 1024)	64226304
dense_7 (Dense)	(None, 512)	524800
dense_8 (Dense)	(None, 4)	2052

```

=====
Total params: 67,011,140
Trainable params: 65,165,316
Non-trainable params: 1,845,824

```

```

-----
None
Do you want model asks you to halt the training [y/n] ?
Epoch      Loss      Accuracy  V_loss    V_acc     LR      Next LR   Monitor  %
Improv      Duration
1 /20      0.994    77.185    0.59761   77.977    0.00100  0.00100   accuracy
0.00      67.89
2 /20      0.448    83.593    0.76103   78.403    0.00100  0.00100   accuracy
8.30      61.26
3 /20      0.371    86.257    0.67452   80.340    0.00100  0.00100   accuracy
3.19      62.03
4 /20      0.328    87.657    0.48797   85.681    0.00100  0.00100   accuracy
1.62      60.86
5 /20      0.305    88.696    0.57162   81.758    0.00100  0.00100   accuracy
1.19      61.87
6 /20      0.295    88.962    0.40935   85.255    0.00100  0.00100   accuracy
0.30      60.80
7 /20      0.272    90.043    0.40576   86.011    0.00100  0.00100   val_loss
0.88      63.08
8 /20      0.258    90.149    0.58749   83.365    0.00100  0.00100   val_loss
-44.79    61.11
9 /20      0.232    91.389    0.38237   86.862    0.00100  0.00100   val_loss
5.77      62.90
10 /20     0.232    91.425    0.40723   87.760    0.00100  0.00100   val_loss
-6.50     64.64

```

```

11 /20      0.226   91.814   0.60903   84.783   0.00100   0.00100   val_loss
-59.28   61.64
12 /20      0.199   93.043   0.46314   84.074   0.00100   0.00050   val_loss
-21.12   61.55
13 /20      0.141   94.933   0.38231   88.658   0.00050   0.00050   val_loss
0.02     61.21
14 /20      0.121   95.576   0.37215   89.083   0.00050   0.00050   val_loss
2.66     61.62
15 /20      0.104   96.368   0.41466   88.422   0.00050   0.00050   val_loss
-11.42   61.64
16 /20      0.093   96.799   0.37416   88.894   0.00050   0.00050   val_loss
-0.54    62.00
17 /20      0.087   96.905   0.45120   88.422   0.00050   0.00025   val_loss
-21.24   62.30
18 /20      0.056   98.210   0.57556   88.705   0.00025   0.00025   val_loss
-54.66   64.13
19 /20      0.049   98.435   0.54046   88.800   0.00025   0.00025   val_loss
-45.22   63.76
20 /20      0.038   98.630   0.58730   89.319   0.00025   0.00013   val_loss
-57.81   61.23

```

training elapsed time was 0.0 hours, 20.0 minutes, 51.50 seconds)

Model: "sequential_3"

Layer (type)	Output Shape	Param #
mobilenetv2_1.00_224 (Functional)	(None, 7, 7, 1280)	2257984
flatten_3 (Flatten)	(None, 62720)	0
dense_9 (Dense)	(None, 1024)	64226304
dense_10 (Dense)	(None, 512)	524800
dense_11 (Dense)	(None, 4)	2052

Total params: 67,011,140

Trainable params: 65,473,156

Non-trainable params: 1,537,984

None

Do you want model asks you to halt the training [y/n] ?

Epoch	Loss	Accuracy	V_loss	V_acc	LR	Next LR	Monitor	%
Improv	Duration							
1 /20	0.945	76.595	10.97996	30.624	0.00100	0.00100	accuracy	
0.00	71.14							
2 /20	0.458	83.404	5.11623	42.580	0.00100	0.00100	accuracy	

8.89	66.15							
3 /20	0.375	86.168	5.98740	34.924	0.00100	0.00100	accuracy	
3.31	66.23							
4 /20	0.340	87.320	4.25912	37.004	0.00100	0.00100	accuracy	
1.34	62.41							
5 /20	0.334	87.609	1.93340	55.766	0.00100	0.00100	accuracy	
0.33	62.56							
6 /20	0.306	88.661	6.28908	25.142	0.00100	0.00100	accuracy	
1.20	62.37							
7 /20	0.301	89.003	2.98812	53.166	0.00100	0.00100	accuracy	
0.39	62.12							
8 /20	0.267	90.019	4.98894	42.864	0.00100	0.00100	val_loss	
-158.04	62.31							
9 /20	0.262	90.385	2.49987	63.469	0.00100	0.00100	val_loss	
-29.30	63.33							
10 /20	0.250	90.804	3.20290	50.095	0.00100	0.00050	val_loss	
-65.66	63.12							
11 /20	0.190	92.919	0.79539	81.474	0.00050	0.00050	val_loss	
58.86	62.78							
12 /20	0.167	93.698	0.60830	85.208	0.00050	0.00050	val_loss	
23.52	61.97							
13 /20	0.156	94.200	0.85854	81.049	0.00050	0.00050	val_loss	
-41.14	60.92							
14 /20	0.142	94.850	0.68134	81.853	0.00050	0.00050	val_loss	
-12.01	62.07							
15 /20	0.137	95.027	0.53209	85.444	0.00050	0.00050	val_loss	
12.53	63.57							
16 /20	0.122	95.582	1.02781	80.246	0.00050	0.00050	val_loss	
-93.16	63.31							
17 /20	0.106	96.208	0.57790	86.153	0.00050	0.00050	val_loss	
-8.61	61.42							
18 /20	0.108	96.244	0.88770	82.278	0.00050	0.00025	val_loss	
-66.83	62.24							
19 /20	0.072	97.525	0.66038	86.200	0.00025	0.00025	val_loss	
-24.11	60.91							
20 /20	0.053	98.228	0.58933	86.626	0.00025	0.00025	val_loss	
-10.76	62.11							

training elapsed time was 0.0 hours, 21.0 minutes, 7.08 seconds)

```
[25]: model1.save(f'saved_model/Model{1}.h5', save_format="h5")
      model2.save(f'saved_model/Model{2}.h5', save_format="h5")
      model3.save(f'saved_model/Model{3}.h5', save_format="h5")
```

```
[27]: !ls /kaggle/working/saved_model
      from IPython.display import FileLink
      display(FileLink("saved_model/Model1.h5"))
      display(FileLink("saved_model/Model2.h5"))
```

```
display(FileLink("saved_model/Model3.h5"))
```

Model1 Model1.h5 Model2 Model2.h5 Model3 Model3.h5


/kaggle/working/saved_model/Model1.h5

/kaggle/working/saved_model/Model2.h5

/kaggle/working/saved_model/Model3.h5

```
[19]: history = [history_3, history_4, history_5 ]
      i = 1
      for his in history:
          print(f'Model_{i}')
          plot_training(his)
          i+=1
```

Model_1



Model_2



Model_3



```
[20]: models = [model1, model2, model3]
i = 1
for model in models:
    ts_length = len(test_df)
    test_batch_size = test_batch_size = max(sorted([ts_length // n for n in
↪range(1, ts_length + 1) if ts_length%n == 0 and ts_length/n <= 80]))
    test_steps = ts_length // test_batch_size

    train_score = model.evaluate(train_gen, steps= test_steps, verbose= 1)
    valid_score = model.evaluate(valid_gen, steps= test_steps, verbose= 1)
    test_score = model.evaluate(test_gen, steps= test_steps, verbose= 1)
    print(f"=====model_{i}=====")
    print("Train Loss: ", train_score[0])
```

```

print("Train Accuracy: ", train_score[1])
print('-' * 20)
print("Validation Loss: ", valid_score[0])
print("Validation Accuracy: ", valid_score[1])
print('-' * 20)
print("Test Loss: ", test_score[0])
print("Test Accuracy: ", test_score[1])
print("=====")
i+=1

```

```

29/29 [=====] - 2s 54ms/step - loss: 0.2619 - accuracy:
0.9052
29/29 [=====] - 2s 50ms/step - loss: 0.3450 - accuracy:
0.8621
29/29 [=====] - 20s 678ms/step - loss: 0.3419 -
accuracy: 0.8781
=====model_1=====
Train Loss: 0.2619060277938843
Train Accuracy: 0.9051724076271057
-----
Validation Loss: 0.3450435996055603
Validation Accuracy: 0.8620689511299133
-----
Test Loss: 0.3418533205986023
Test Accuracy: 0.8781294226646423
=====
29/29 [=====] - 2s 54ms/step - loss: 0.0881 - accuracy:
0.9698
29/29 [=====] - 2s 52ms/step - loss: 0.4603 - accuracy:
0.8728
29/29 [=====] - 8s 282ms/step - loss: 0.3341 -
accuracy: 0.9017
=====model_2=====
Train Loss: 0.08808482438325882
Train Accuracy: 0.9698275923728943
-----
Validation Loss: 0.46028727293014526
Validation Accuracy: 0.8728448152542114
-----
Test Loss: 0.334089070558548
Test Accuracy: 0.9017477631568909
=====
29/29 [=====] - 2s 54ms/step - loss: 0.2328 - accuracy:
0.9203
29/29 [=====] - 1s 47ms/step - loss: 0.5803 - accuracy:
0.8384
29/29 [=====] - 8s 267ms/step - loss: 0.5104 -

```

```

accuracy: 0.8715
=====model_3=====
Train Loss:  0.23283062875270844
Train Accuracy:  0.920258641242981
-----
Validation Loss:  0.580291748046875
Validation Accuracy:  0.8383620977401733
-----
Test Loss:  0.510432243347168
Test Accuracy:  0.8715162873268127
=====

```

```

[28]: Preds = {}
      i = 1
      for model in models:
          preds = model.predict_generator(test_gen)
          y_pred = np.argmax(preds, axis=1)
          Preds[f"Model{i}"] = y_pred
          i+=1

      print(Preds)

```

```

{'Model1': array([1, 1, 1, ..., 2, 2, 2]), 'Model2': array([1, 1, 1, ..., 2, 2,
2]), 'Model3': array([1, 1, 1, ..., 2, 2, 2])}

```

```

[36]: i =1
      for model in models:
          print(f"=====Model {i}=====")
          y_pred = Preds[f"Model{i}"]
          g_dict = test_gen.class_indices
          classes = list(g_dict.keys())

          # Confusion matrix
          cm = confusion_matrix(test_gen.classes, y_pred)
          plot_confusion_matrix(cm= cm, classes= classes, title = f'Confusion Matrix_
          ↪model {i}')

          # Classification report
          print(classification_report(test_gen.classes, y_pred, target_names=
          ↪classes))
          i+=1
          print("=====")

```

```

=====Model 1=====
Confusion Matrix, Without Normalization
[[300  17  42   3]
 [  9 481 112   0]
 [ 11  53 952   3]

```

```

[ 1  0  7 126]]
      precision    recall  f1-score   support

      COVID          0.93      0.83      0.88        362
    Lung_Opacity      0.87      0.80      0.83        602
        Normal      0.86      0.93      0.89       1019
Viral Pneumonia      0.95      0.94      0.95        134

    accuracy              0.88       2117
   macro avg          0.90      0.88      0.89       2117
  weighted avg          0.88      0.88      0.88       2117

```

=====

=====Model 2=====

Confusion Matrix, Without Normalization

```

[[331  13  15   3]
 [  6 505  91   0]
 [  9  65 944   1]
 [  0   0   5 129]]
      precision    recall  f1-score   support

      COVID          0.96      0.91      0.94        362
    Lung_Opacity      0.87      0.84      0.85        602
        Normal      0.89      0.93      0.91       1019
Viral Pneumonia      0.97      0.96      0.97        134

    accuracy              0.90       2117
   macro avg          0.92      0.91      0.92       2117
  weighted avg          0.90      0.90      0.90       2117

```

=====

=====Model 3=====

Confusion Matrix, Without Normalization

```

[[338   8  14   2]
 [ 14 512  76   0]
 [ 50  98 869   2]
 [  1   1   6 126]]
      precision    recall  f1-score   support

      COVID          0.84      0.93      0.88        362
    Lung_Opacity      0.83      0.85      0.84        602
        Normal      0.90      0.85      0.88       1019
Viral Pneumonia      0.97      0.94      0.95        134

    accuracy              0.87       2117
   macro avg          0.88      0.89      0.89       2117
  weighted avg          0.87      0.87      0.87       2117

```

=====

